

Introduction to Deep Learning

Ups and Downs of Deep Learning

- 1958, perceptron (linear model)
- 1969: Perceptron has limitation
- 1980: Multi-layer perceptron (X)
 - Do not have significant difference from DNN today
- 1986: Backpropagation
 - Usually more than 3 hidden layers is not helpful
- 1989: 1 hidden layer is 'good enough', why deep learning?
- 2006: RBM initialization (breakthrough)
- 2009: GPU
- 2011: Start to be popular in speech recognition
- 2012: win ILSVRC image competition

Three Steps for Deep Learning

- Define a set of function (Neural Network)
- Evaluate the performance of function
- Pick up the best function

Step1:

Neural Network


Concatenate different logistic regressions, each logistic regression has its own weight and bias.

Different connection leads to different network structures:

Network parameter theta: all the weights and biases in the 'neurons'


How to connect them?

Fully Connect Feedforward Network



Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Created with EverCam
<http://www.camdemy.com>

weight and bias are from training data.


for example, blue top neuron, weight (1, -2), bias(1); blue bottom neuron, weight (-1, 1), bias(0)

top blue output: $1 \times 1 + (-1) \times (-2) + 1 \text{ (bias)} = 4 \gg \gg \text{Pass the sigmoid function} \gg \gg 0.98$

bottom blue output: $1 \times (-1) + (-1) \times 1 + 0 = -2 \gg \gg \text{Pass the sigmoid function} \gg \gg 0.12$

Continuing ...

Fully Connect Feedforward Network




Input is a vector (1, -1), output is a vector (0.62, 0.83)

Given network structure, define a function set

Types of connections


Fully Connect Feedforward Network



Deep = many hidden layers

Matrix Operation

Matrix Operation



$$\sigma \left(\underbrace{\begin{bmatrix} 1 & -2 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}}_{\begin{bmatrix} 4 \\ -2 \end{bmatrix}} \right) = \begin{bmatrix} 0.98 \\ 0.12 \end{bmatrix}$$

Created with EverCam.
<http://www.camdemy.co>


Feedforward Network Calculation Process:

(input multiply by matrix (weight) + bias) >> into sigmoid (any activation function) >>> output

$$\begin{aligned} &\sigma(W^1 * x + b^1) \\ &\sigma(W^2 * a^1 + b^2) \\ &\dots \end{aligned}$$

a series of matrix computations

Output Layer




output layer = multi-class classifier

output layer: after the hidden layers conducted complex transformation, it (output layer) takes a set of features

Example Application




Input



Ink \rightarrow 1

No ink \rightarrow 0

Output



Each dimension represents the confidence of a digit.

Created with EverCam.

Softmax >>> y is a 10-dim vector, each dimension is corresponding to a number

The only constraint for this hand-written image recognition is the 256-dim input, and 10-dim output.

The number of layers and how they connected is unlimited, you can design this part.

Q: How many layers? How many neurons for each layer?

A: Trail and Error + Intuition

Q: Can the structure be automatically determined?

A: e.g. Evolutionary Artificial Neural Networks

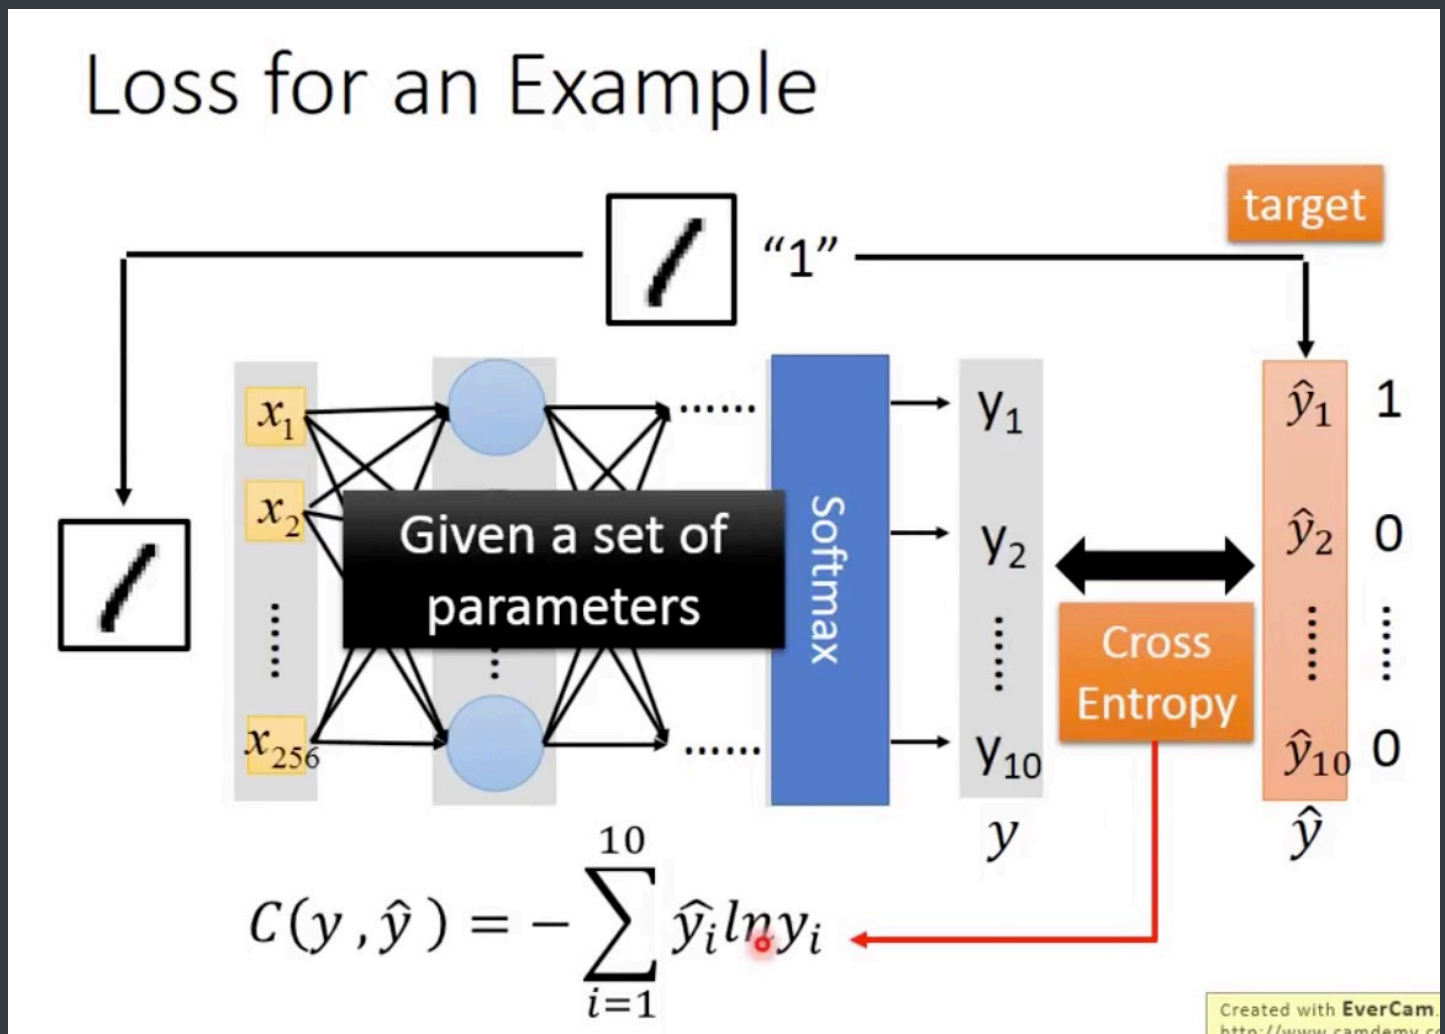
Q: Can we design the network structure?

A: CNN

Note: From ML to DL, the question transformed from 'Feature Extraction' to 'Network Design'


Step2:

Loss function

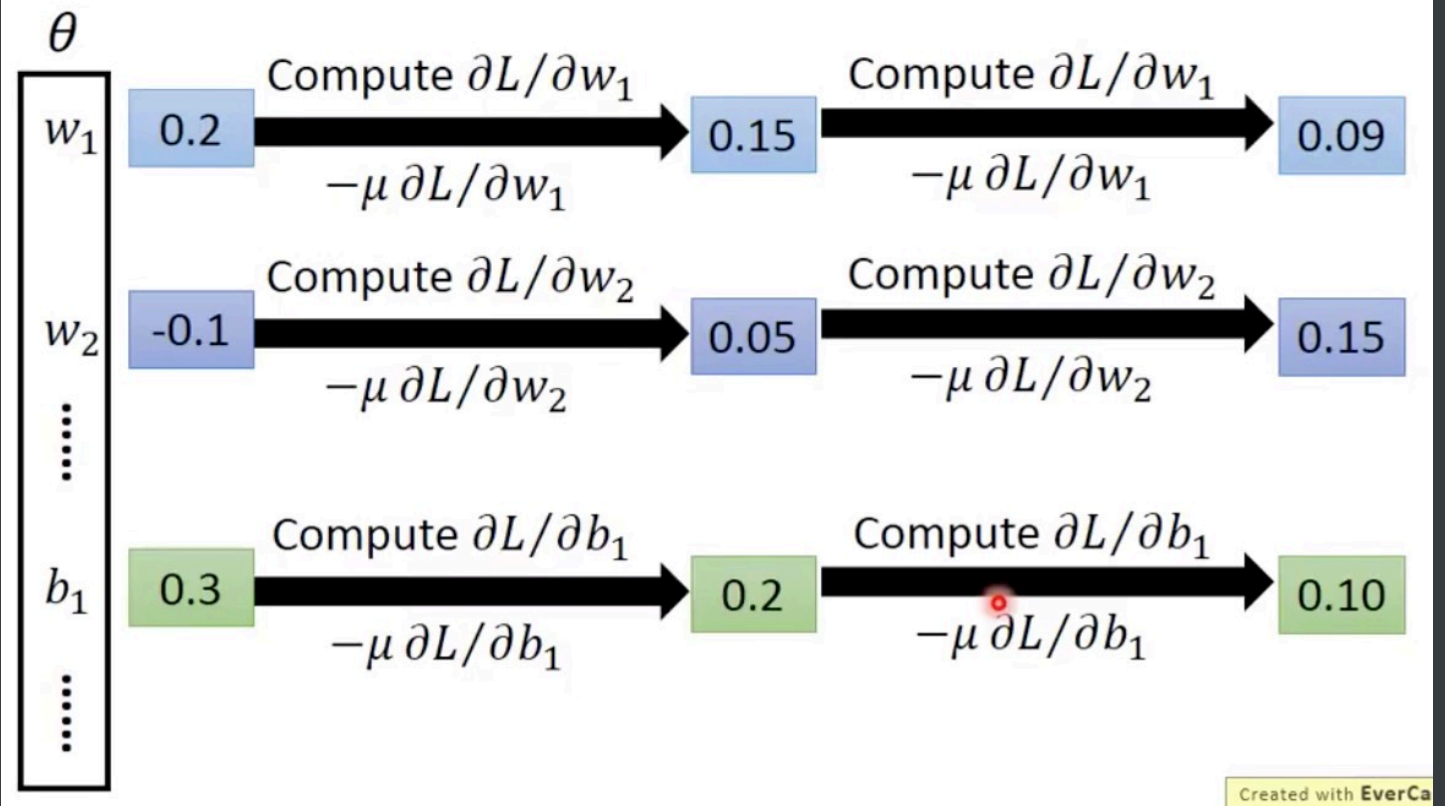


Calculate the cross entropy of y and \hat{y} , find the parameter that minimizes the loss.

Gradient Descent



Gradient Descent



Backpropagation

An efficient way to compute

$$\frac{dL}{dW}$$

Packages:

TensorFlow, Torch, Theano, Caffe, CNTK, Chainer, DSSTNE, mxnet, liban

Resources on DL:

- ML and having it deep and structured
- Neural Networks and Deep Learning: Michael Nielsen
- Deep Learning: Yoshua Bengio, Ian J. Goodfellow and Aaron Courville

